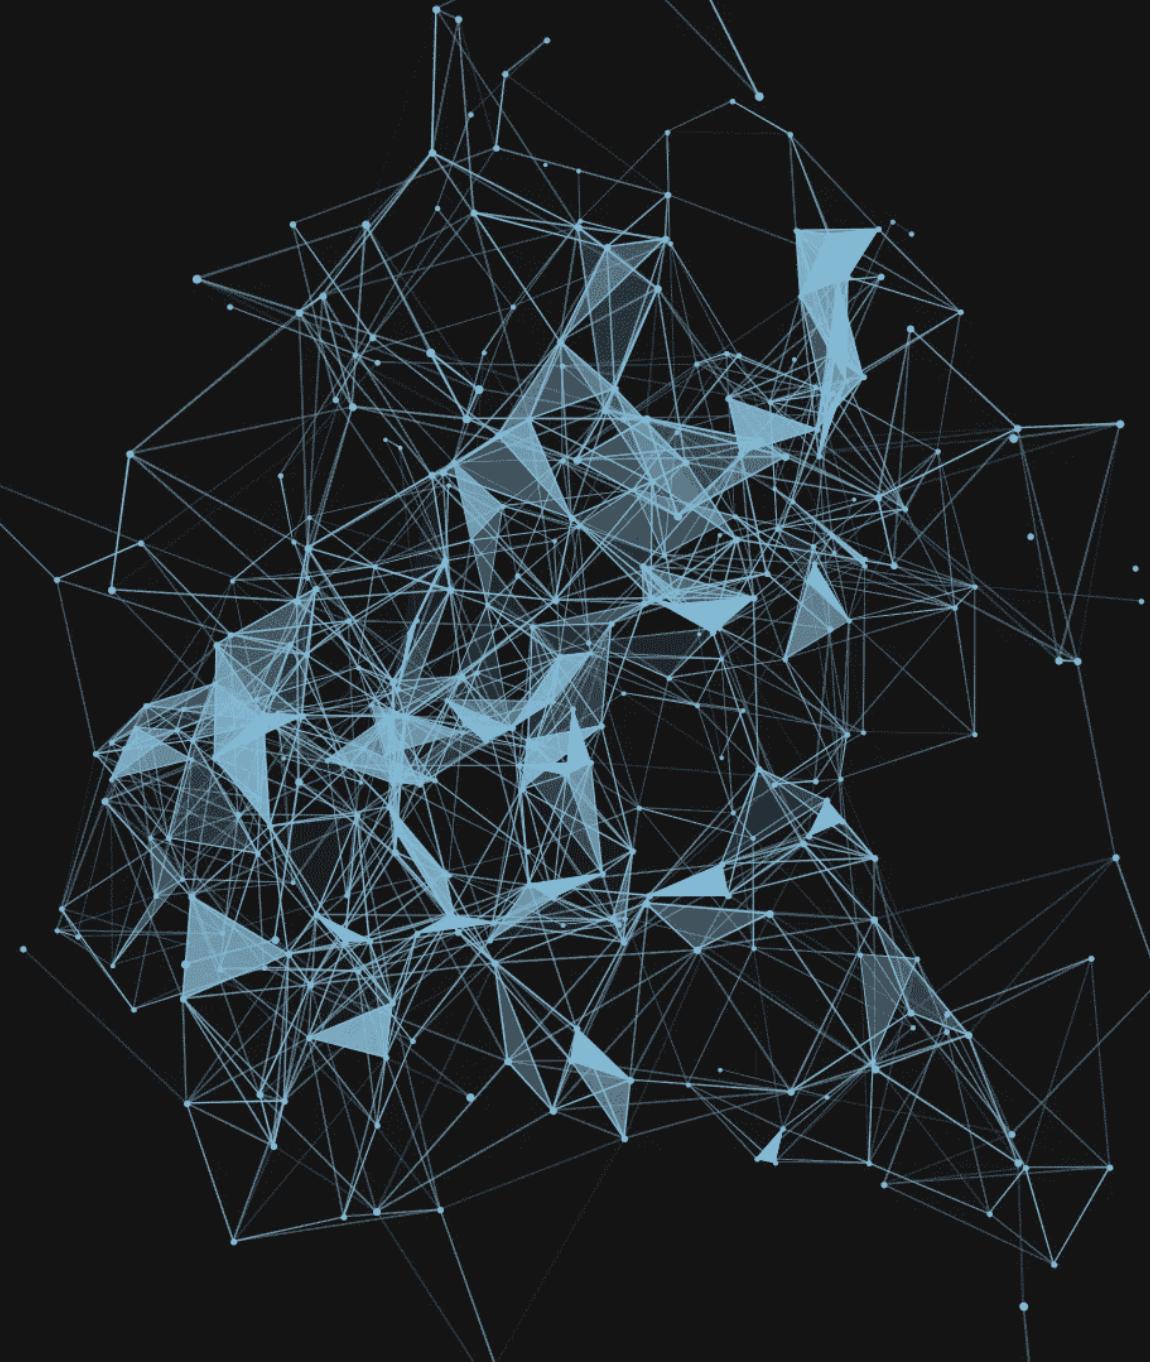


R + Analytics + Cloud

Analytics/Big Data Training



Thiago Pires | 22 and 23 Apr 2021
github.com/th1460/r-analytics



Analytics

Fact-based decisions have become our competitive strength. Using or not using analytics is no longer an option

Analytics

Analytics is a comprehensive and multidimensional field that uses mathematical techniques, statistics, predictive modeling and machine learning to find meaningful patterns and knowledge in data.

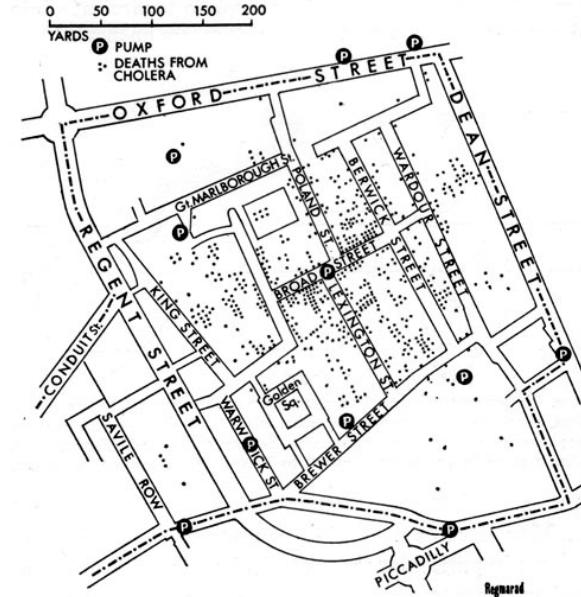
Florence Nightingale recording and analyzing mortality data in the 1850s. She presented data on wounded soldiers with death counts per month. He ended up realizing that the main **death rate was due to hospital infection and not battle injuries**, as many imagined.



Analytics

A concrete example of spatial analysis carried out in 1854 in the city of London, where the population was suffering a serious cholera epidemic, disease about which at the time the form of contamination was not known. In a situation where more than 500 deaths had already occurred, the Dr. John Snow had an idea: put on the city map the location of the cholera patients and the water wells (at that time, the main source of water for the city's inhabitants). The obtained map is shown in the figure below.

Dr. Snow realized that most cases were concentrated around the "Broad Street" well



Analytics

Data analysis can reveal correlations and patterns. There is less need to rely on assumptions or intuition. And it can help answer questions like:

- What happened?
- How or why it happened?
- What is happening now?
- What will likely happen next?

R

Linguagem R

Developed by **Ross Ihaka** and **Robert Gentleman** in the Statistics department at the University of Auckland, New Zealand in 1993.



Exemplo

```
1 + 2 + (3 * 4)
```

```
[1] 15
```

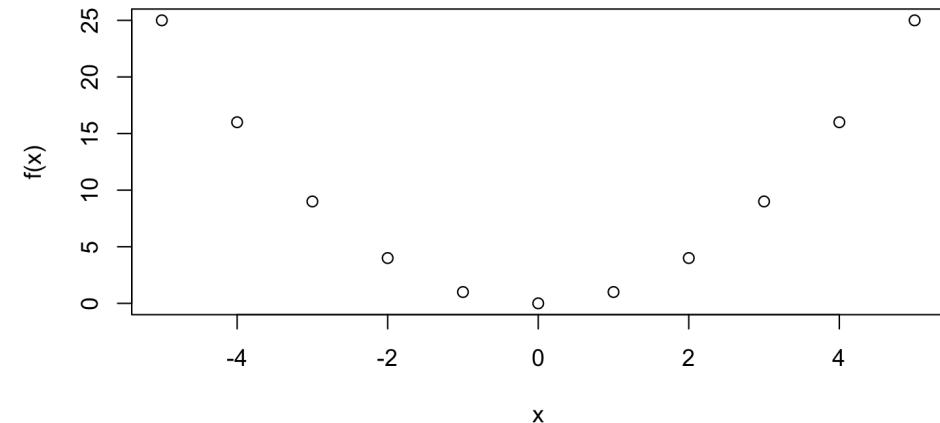
```
f <- function(x) {x^2}
x <- seq(-5, 5, 1); f(x)
```

```
[1] 25 16 9 4 1 0 1 4 9 16 25
```

```
paste("If x = 2, then f(2) =", f(2))
```

```
[1] "If x = 2, then f(2) = 4"
```

```
plot(x, f(x))
```



Base language × Tidyverse

Base language

```
# sum elements in a vector  
sum(c(10, 5, 8, 12, 2, 0))
```

```
[1] 37
```

```
# data.frame  
df <- data.frame(x = 1:3, y = letters[1:3])  
df$z <- df$x^2
```

```
# filter and select  
df[df$y == "b", c("x", "z")]
```

```
x z  
2 2 4
```

Tidyverse

```
require(dplyr)  
# sum elements in a vector  
c(10, 5, 8, 12, 2, 0) %>% sum
```

```
[1] 37
```

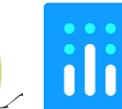
```
# tibble  
df <- tibble(x = 1:3, y = letters[1:3], z = x^2)
```

```
# filter and select  
df %>% filter(y == "b") %>% select(x, z)
```

```
# A tibble: 1 x 2  
      x     z  
   <int> <dbl>  
1     2     4
```

Packages



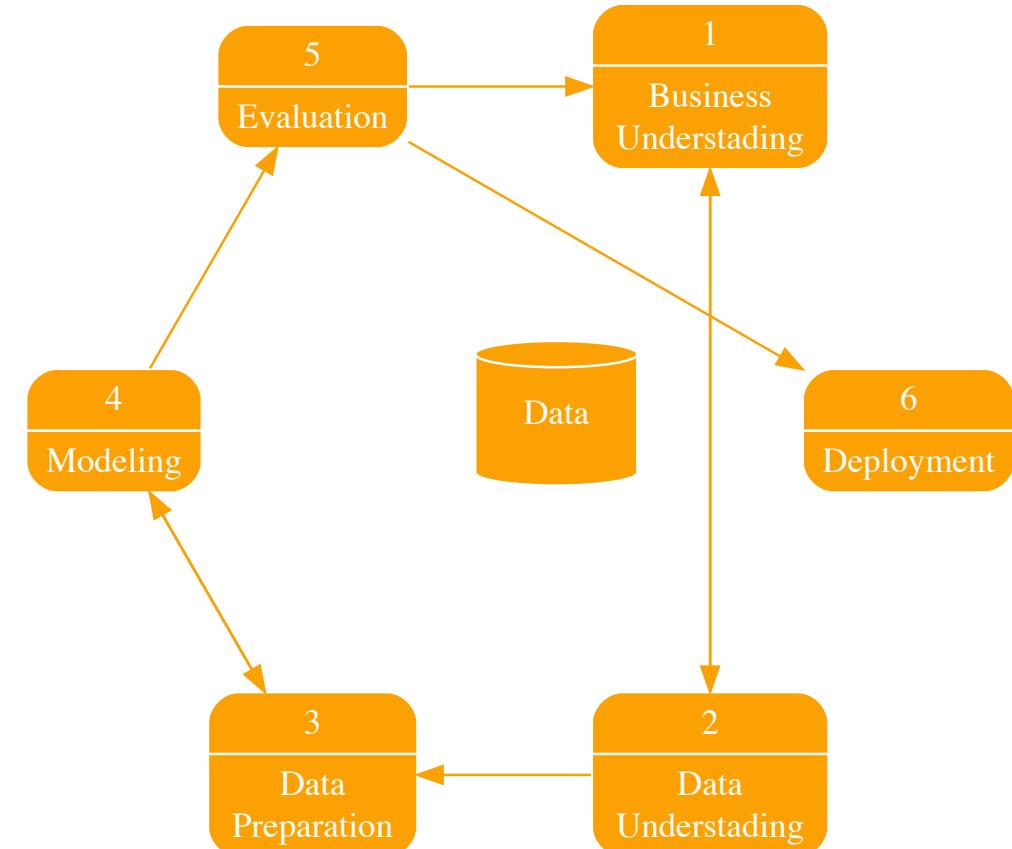
Leaflet  plotly 

CRISP-DM

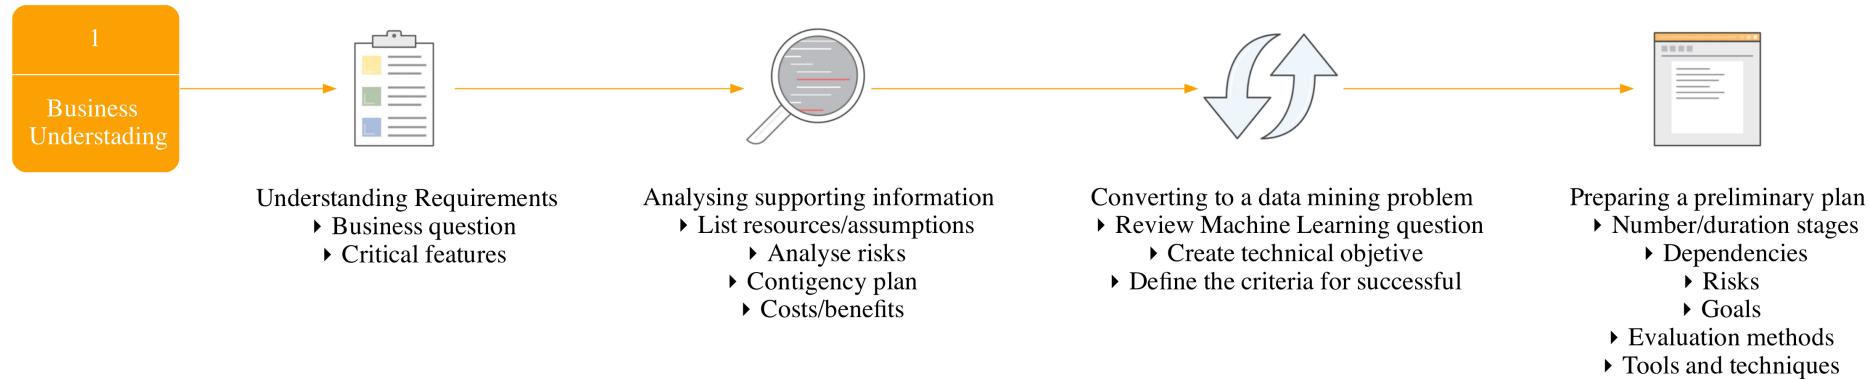
| *Cross Industry Standard Process for Data Mining*

CRISP-DM

Cross-industry standard process for data mining, known as CRISP-DM is an open standard process model that describes common approaches used by data mining experts. It is the most widely-used analytics model



Business Understanding



Case Titanic dataset challenge

This is a *part* (train data) from the dataset used in Kaggle Titanic challenge

Business question:

What factors was associated with a person survive in the Titanic disaster?

Features:

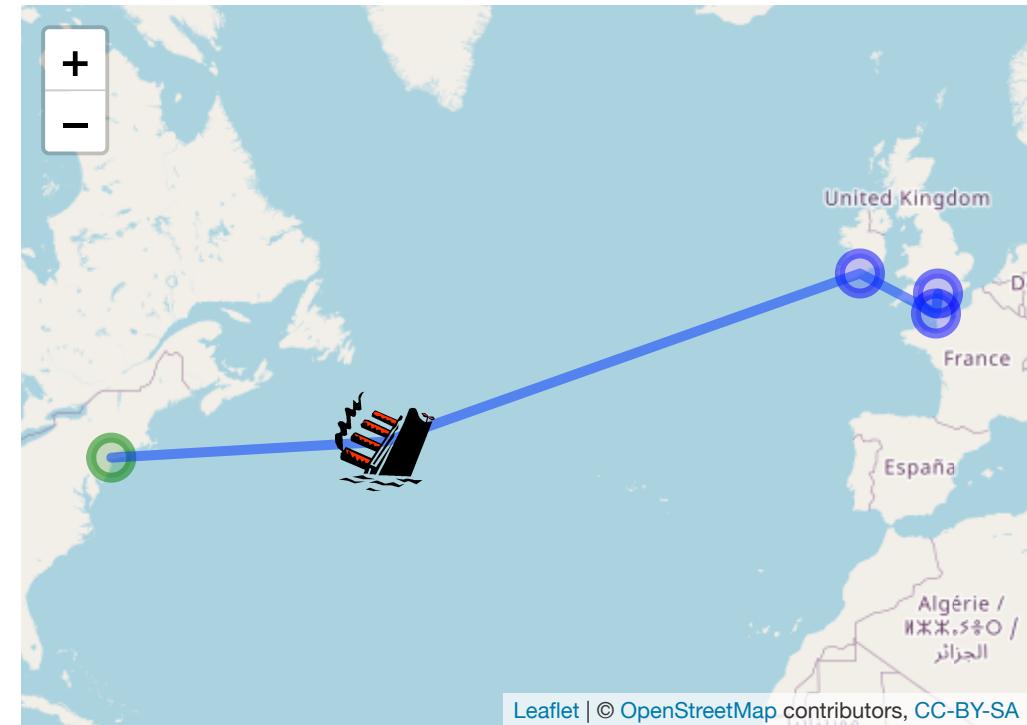
- How to analyze the data? (e.g. R, Python, etc)
- How prepare the data? (e.g. Pandas, dplyr, etc)
- What were the passenger types (e.g. Ages, Gender, Class, etc)

Titanic disaster

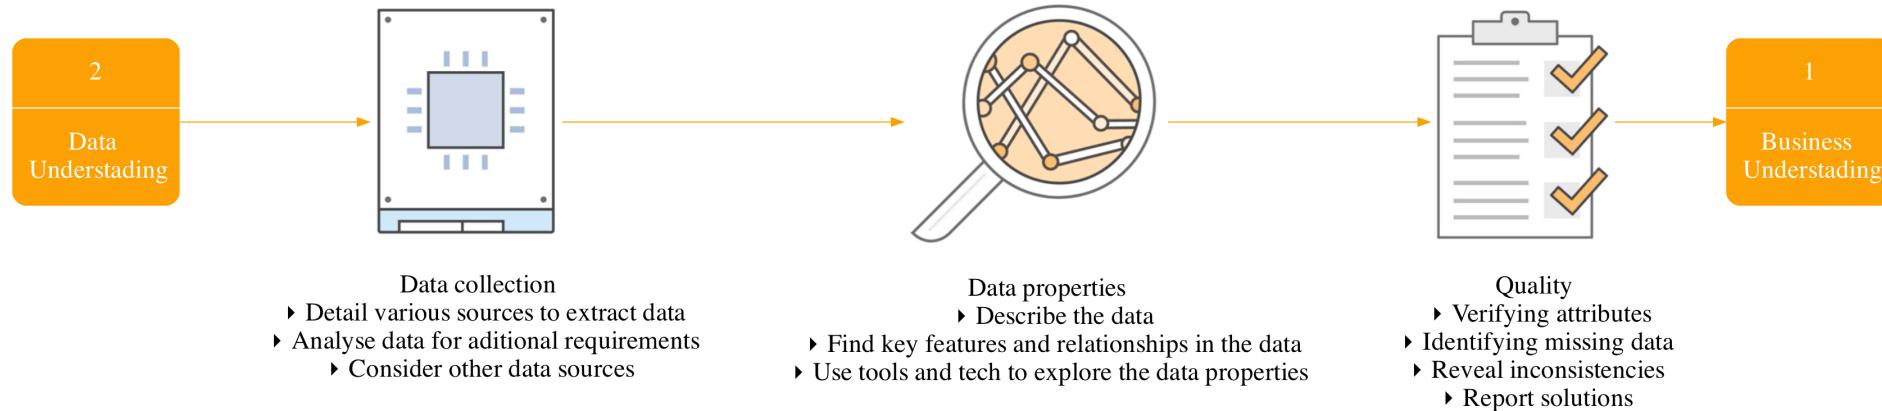
```
require(leaflet)

events <- bind_rows(
  tibble(location = "Southampton (10-04-1912)",
         lng = -1.4191, lat = 50.7894),
  tibble(location = "Cherbourg (10-04-1912)",
         lng = -1.6109, lat = 49.6445),
  tibble(location = "Queenstown (11-04-1912)",
         lng = -8.3211, lat = 51.8535),
  tibble(location = "Naufrágio (14-04-1912)",
         lng = -49.9408, lat = 41.7258),
  tibble(location = "New York",
         lng = -73.9655, lat = 40.6832))

leaflet() %>% addTiles() %>%
  addCircleMarkers(data = events %>% slice(1:3, 5),
                  label = ~location,
                  color = c(rep("blue", 3), "green"))
  addMarkers(data = events %>% slice(4),
             icon = list(
               iconUrl = "resources/images/sinking-shi
               iconSize = c(50, 50)),
             label = ~location) %>%
  addPolylines(data = events, ~lng, ~lat)
```



Data Understanding



Titanic case

```
# Read data from Github
dataset <- "https://gist.githubusercontent.com/michhar/2dfd2de0d4f8727f873422c5d959fff5/raw/fa71405126017e6a37bea592"
```

Data Understanding (Titanic case)

Study the data dictionary and

Inspect the dataset

```
titanic <- readr::read_csv(dataset)
titanic %>% glimpse()
```

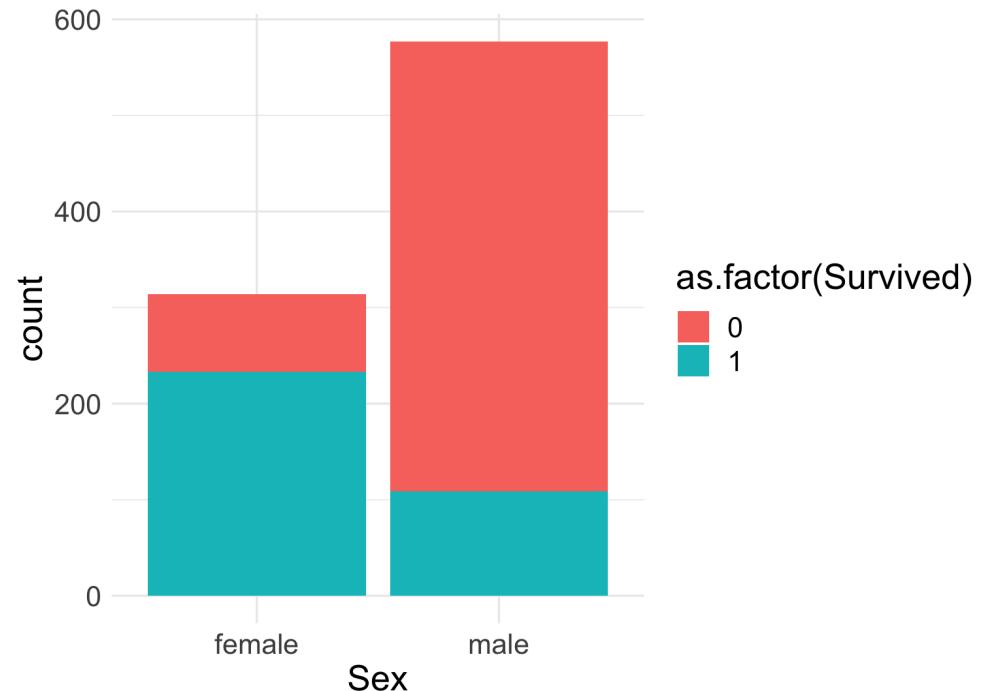
```
Rows: 891
Columns: 12
$ PassengerId <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, ...
$ Survived      <dbl> 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1...
$ Pclass        <dbl> 3, 1, 3, 1, 3, 3, 1, 3, 3, 2, 3, 1, 3, 3, 3, 2, 3, 2, 3, 3...
$ Name          <chr> "Braund, Mr. Owen Harris", "Cumings, Mrs. John Bradley (Fl...
$ Sex           <chr> "male", "female", "female", "female", "male", "male", "mal...
$ Age           <dbl> 22, 38, 26, 35, 35, NA, 54, 2, 27, 14, 4, 58, 20, 39, 14, ...
$ SibSp         <dbl> 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 0, 0, 1, 0, 0, 4, 0, 1, 0...
$ Parch         <dbl> 0, 0, 0, 0, 0, 0, 1, 2, 0, 1, 0, 0, 5, 0, 0, 1, 0, 0, 0...
$ Ticket        <chr> "A/5 21171", "PC 17599", "STON/O2. 3101282", "113803", "37...
$ Fare          <dbl> 7.2500, 71.2833, 7.9250, 53.1000, 8.0500, 8.4583, 51.8625, ...
$ Cabin         <chr> NA, "C85", NA, "C123", NA, NA, "E46", NA, NA, NA, "G6", "C...
$ Embarked      <chr> "S", "C", "S", "S", "S", "Q", "S", "S", "S", "C", "S", "S"...
```

Data Understanding (Titanic case)

Detect key associations

```
require(ggplot2)
theme_set(
  theme_minimal() +
    theme(text = element_text(size = 18))
)

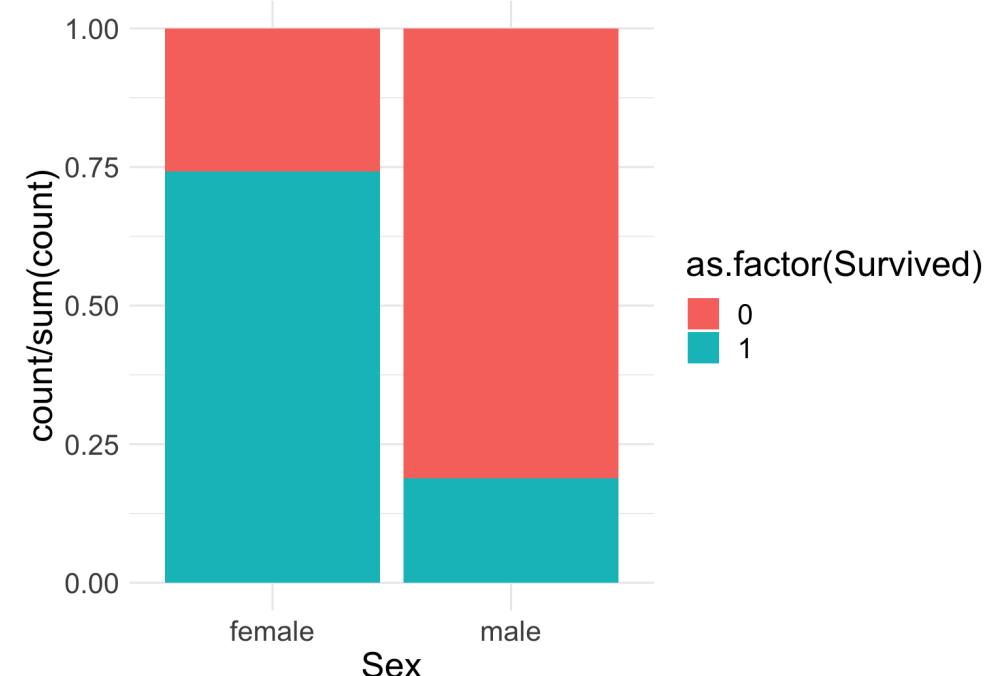
titanic %>%
  ggplot(aes(Sex, ..count..,
             group = Survived,
             fill = as.factor(Survived))) +
  geom_bar()
```



Data Understanding (Titanic case)

Detect key associations

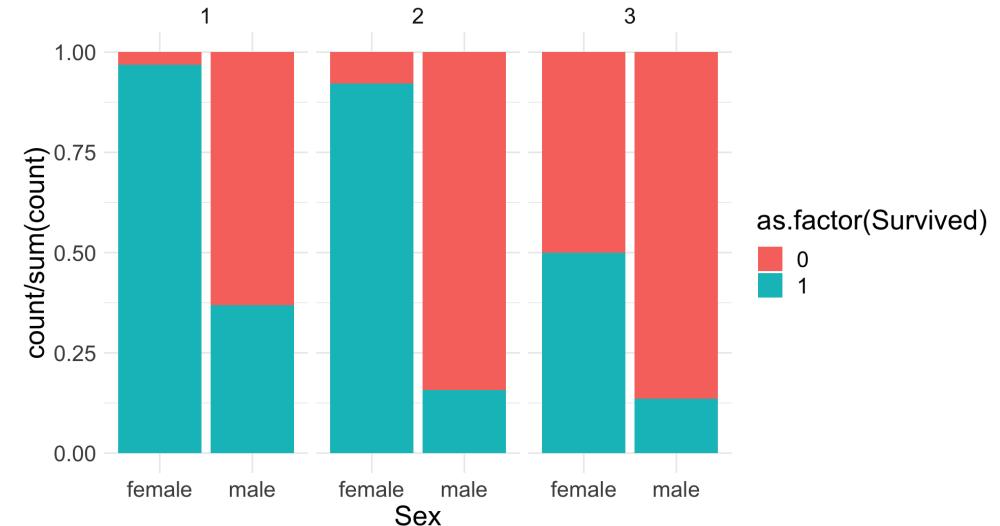
```
titanic %>%  
  ggplot(aes(Sex, ..count../sum(..count..),  
             group = Survived,  
             fill = as.factor(Survived))) +  
  geom_bar(position="fill")
```



Data Understanding (Titanic case)

Detect key associations

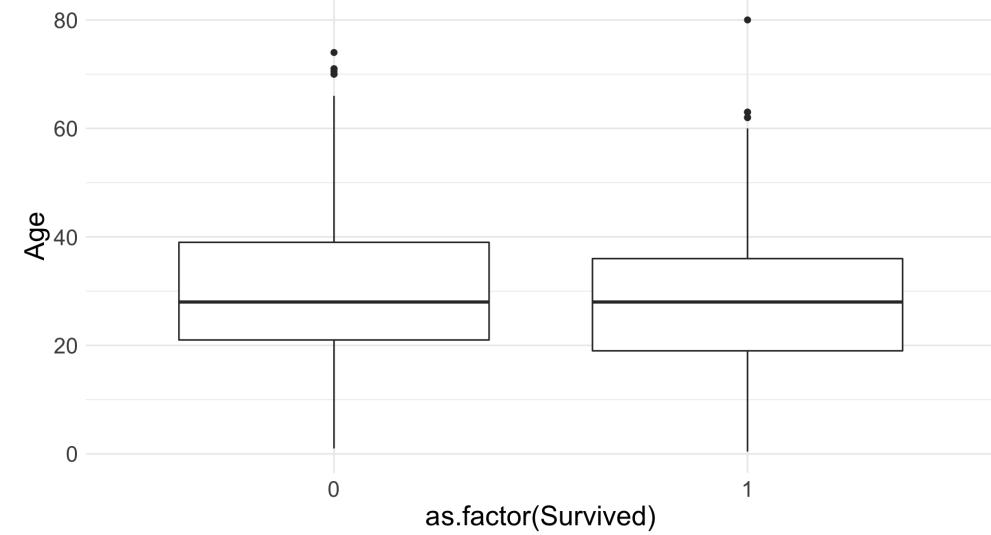
```
titanic %>%
  ggplot(aes(Sex, ..count../sum(..count..),
             group = Survived,
             fill = as.factor(Survived))) +
  geom_bar(position="fill") +
  facet_grid(~Pclass)
```



Data Understanding (Titanic case)

Detect key associations

```
titanic %>%
  ggplot(aes(as.factor(Survived), Age)) +
  geom_boxplot()
```



Data Understanding (Titanic case)

Identify missing data

```
na_count <- function(x) sum(is.na(x))

titanic %>%
  summarise(across(everything(),
    list(na_count),
    .names = "{.col}")) %>%
  tidyr::pivot_longer(everything(),
    values_to = "n_missing")
```

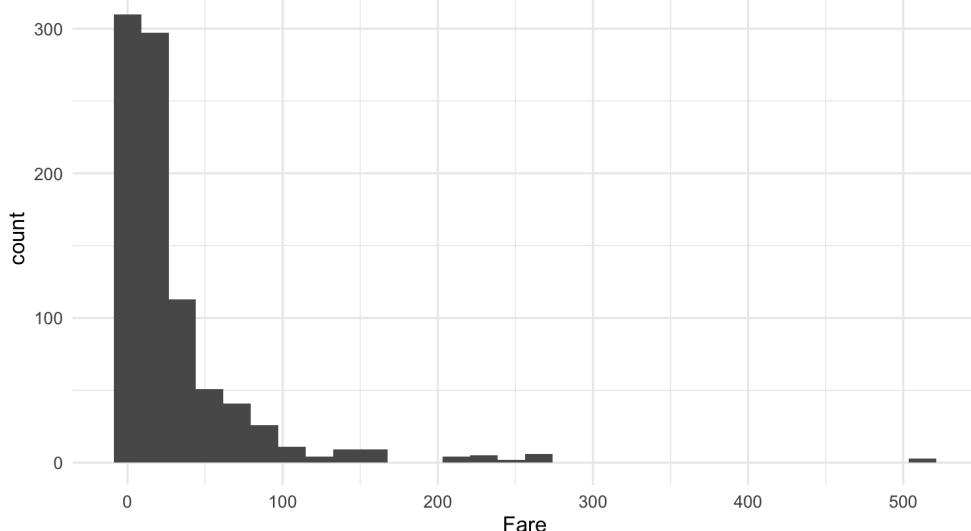
| | # A tibble: 12 x 2 |
|----|--------------------|
| | name n_missing |
| | <chr> <int> |
| 1 | PassengerId 0 |
| 2 | Survived 0 |
| 3 | Pclass 0 |
| 4 | Name 0 |
| 5 | Sex 0 |
| 6 | Age 177 |
| 7 | SibSp 0 |
| 8 | Parch 0 |
| 9 | Ticket 0 |
| 10 | Fare 0 |
| 11 | Cabin 687 |
| 12 | Embarked 2 |

Data Understanding (Titanic case)

Analyse distribution

```
require(ggplot2)

titanic %>%
  ggplot(aes(x = Fare)) +
  geom_histogram() +
  theme_minimal()
```



```
titanic %>%
  summarise(
    min(Fare),
    max(Fare),
    mean(Fare),
    median(Fare),
    e1071::skewness(Fare)
  ) %>%
  tidyverse::pivot_longer(everything(),
                         names_to = "measure") %>%
  mutate(value = value %>% round(2))
```

| measure | value |
|-------------------------|-------|
| <chr> | <dbl> |
| 1 min(Fare) | 0 |
| 2 max(Fare) | 512. |
| 3 mean(Fare) | 32.2 |
| 4 median(Fare) | 14.4 |
| 5 e1071::skewness(Fare) | 4.77 |

Data Understanding (Titanic case)

Symetric distribution

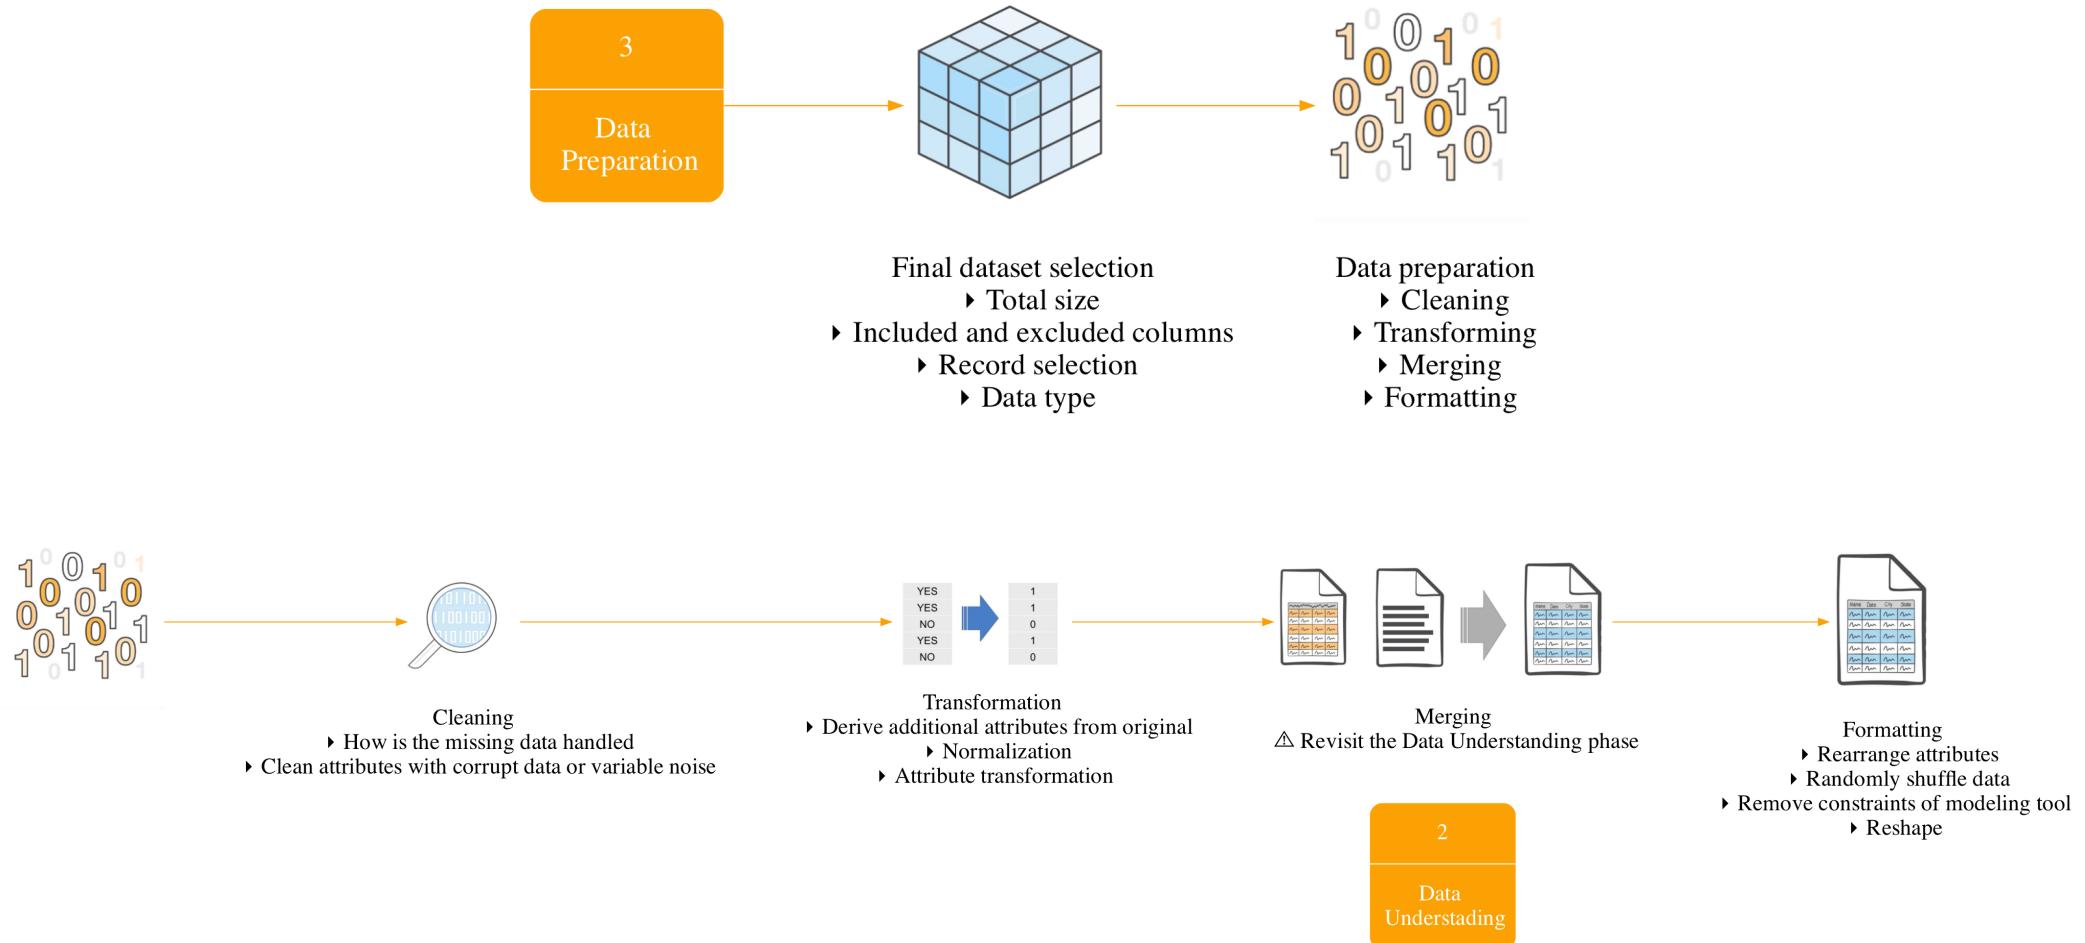
```
set.seed(123)

dataset <- tibble(x = rnorm(100, 30, 2))
dataset %>%
  ggplot(aes(x)) +
  geom_histogram(aes(y = ..density..)) +
  geom_density(aes(x), color = "blue") +
  theme_minimal()
```

```
dataset %>%
  summarise(
    min(x),
    max(x),
    mean(x),
    median(x),
    e1071::skewness(x)
  ) %>%
  tidyverse::pivot_longer(everything(),
                         names_to = "measure") %>%
  mutate(value = value %>% round(2))
```

```
# A tibble: 5 x 2
  measure      value
  <chr>       <dbl>
1 min(x)     25.4
2 max(x)     34.4
3 mean(x)    30.2
4 median(x)  30.1
5 e1071::skewness(x) 0.06
```

Data Preparation



Data Preparation (Titanic case)

Do we drop rows?
Which?

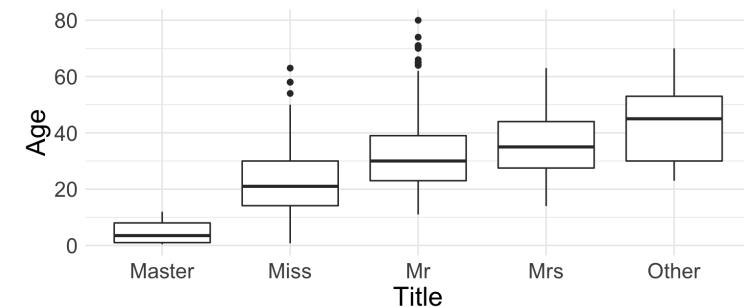
How treat empty
values?

```
# A tibble: 12 x 2
  name      n_missing
  <chr>        <int>
1 PassengerId     0
2 Survived        0
3 Pclass          0
4 Name            0
5 Sex             0
6 Age           177
7 SibSp          0
8 Parch          0
9 Ticket          0
10 Fare           0
11 Cabin         687
12 Embarked       2
```

```
titanic$name[4]
```

```
[1] "Futrelle, Mrs. Jacques Heath (Lily May Peel)"
```

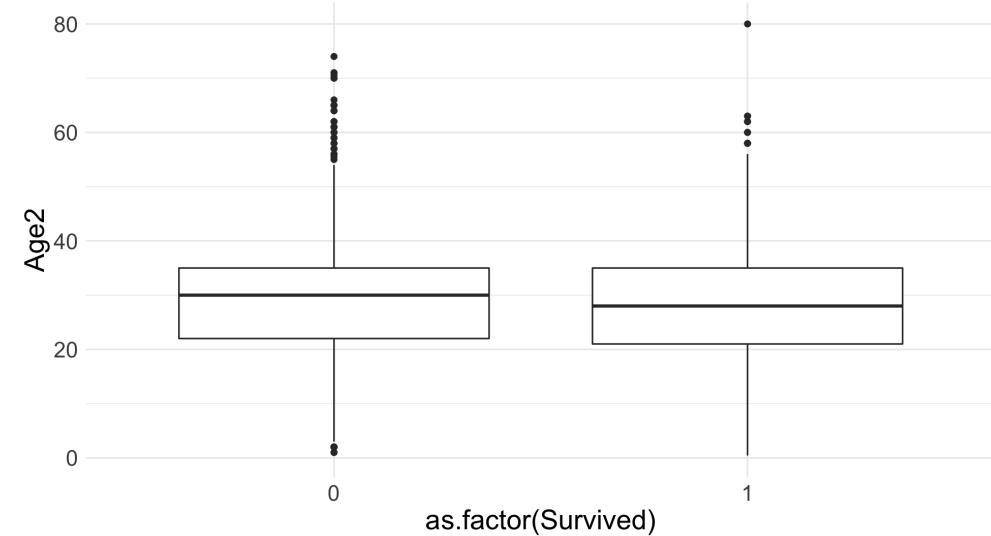
```
require(forcats)
require(ggplot2)
titanic %>%
  mutate(
    Title = Name %>%
      stringr::str_extract(
        "(?=<\\,\\s)(.*)(?=\\.)") %>%
      fct_lump(n = 4)) %>%
  ggplot(aes>Title, Age)) + geom_boxplot
```



Data Preparation (Titanic case)

After fix empty ages

```
titanic %>%
  mutate(
    Title = Name %>%
      stringr::str_extract("(?<=\\,\\s)(.*)(?=\\.\\.)") %
      fct_lump(n = 4)) %>%
  group_by(Title) %>%
  mutate(Age2 = ifelse(is.na(Age), median(Age, na.rm =
  ggplot(aes(as.factor(Survived), Age2)) +
  geom_boxplot()
```



Data Preparation (Titanic case)

Fix types and labels

```
require(magrittr)
titanic %<%>
  mutate(Sex = as.factor(Sex),
         Pclass = factor(Pclass,
                           labels = c("1st", "2nd", "3rd")),
         Survived = factor(Survived,
                           labels = c("No", "Yes")),
         Embarked = fct_recode(Embarked,
                               "Cherbourg" = "C",
                               "Queenstown" = "Q",
                               "Southampton" = "S"))
titanic %>% select(Sex, Survived, Embarked) %>%
  print(n = 4)
```

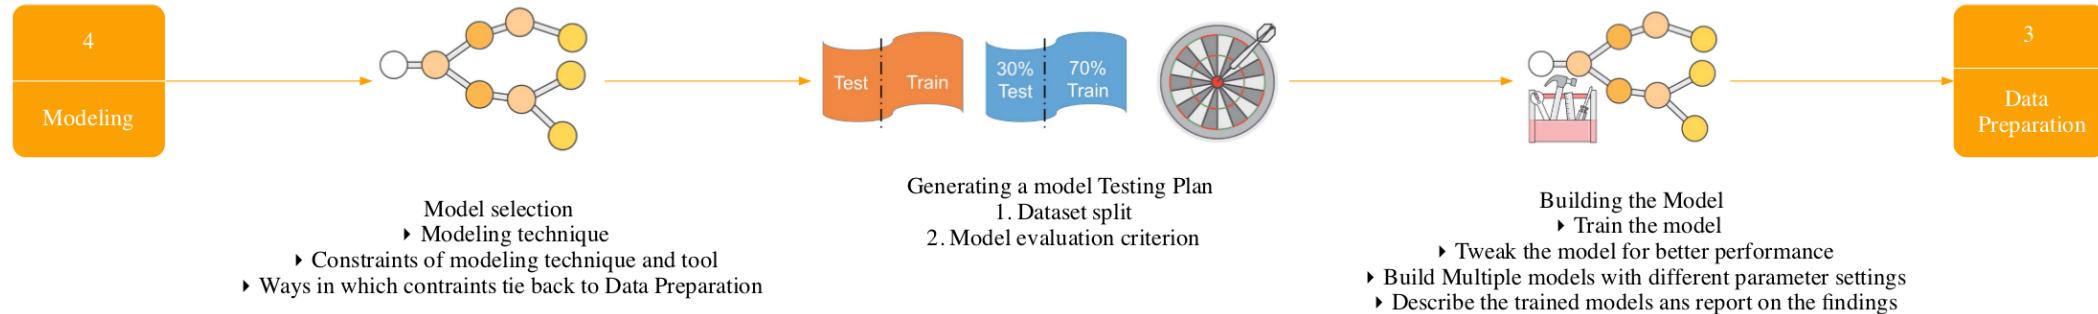
```
# A tibble: 891 x 3
  Sex   Survived Embarked
  <fct> <fct>   <fct>
1 male   No      Southampton
2 female Yes     Cherbourg
3 female Yes     Southampton
4 female Yes     Southampton
# ... with 887 more rows
```

```
freq <- function(data, x) {
  table <- data %>%
    count({{x}}) %>%
    mutate(`%` = (n/sum(n) * 100) %>% round(1)) %>%
    rename(Levels = 1, N = 2)
  tibble(Variable = quo_name(quo({{x}}))) %>%
    bind_cols(table)
}

c("Sex", "Survived", "Embarked", "Pclass") %>%
  purrr::map_dfr(~ freq(titanic, !! sym(.x)))
```

| | Variable | Levels | N | `%` |
|----|----------|-------------|-------|-------|
| | <chr> | <fct> | <int> | <dbl> |
| 1 | Sex | female | 314 | 35.2 |
| 2 | Sex | male | 577 | 64.8 |
| 3 | Survived | No | 549 | 61.6 |
| 4 | Survived | Yes | 342 | 38.4 |
| 5 | Embarked | Cherbourg | 168 | 18.9 |
| 6 | Embarked | Queenstown | 77 | 8.6 |
| 7 | Embarked | Southampton | 644 | 72.3 |
| 8 | Embarked | NA | 2 | 0.2 |
| 9 | Pclass | 1st | 216 | 24.2 |
| 10 | Pclass | 2nd | 184 | 20.7 |
| 11 | Pclass | 3rd | 491 | 55.1 |

Modeling



Titanic case

What factors was associated with a person survive in the Titanic disaster?

Model selection

- Logistic regression, Decision Tree, etc
- Features: Sex, Pclass, interaction between Sex and Pclass?
- Split strategy

Modeling (Titanic case)

Split data

```
require(tidymodels)
set.seed(555)

# Put 3/4 of the data into the training set
data_split <- initial_split(titanic, prop = 3/4)

# Create data frames for the two sets:
train_data <- training(data_split)
test_data <- testing(data_split)
```

Modeling (Titanic case)

Survived ~ Sex + Pclass

```
lr_mod <-
  logistic_reg() %>%
  set_engine("glm")

lr_fit1 <-
  lr_mod %>%
  fit(Survived ~ Sex + Pclass, data = train_data)

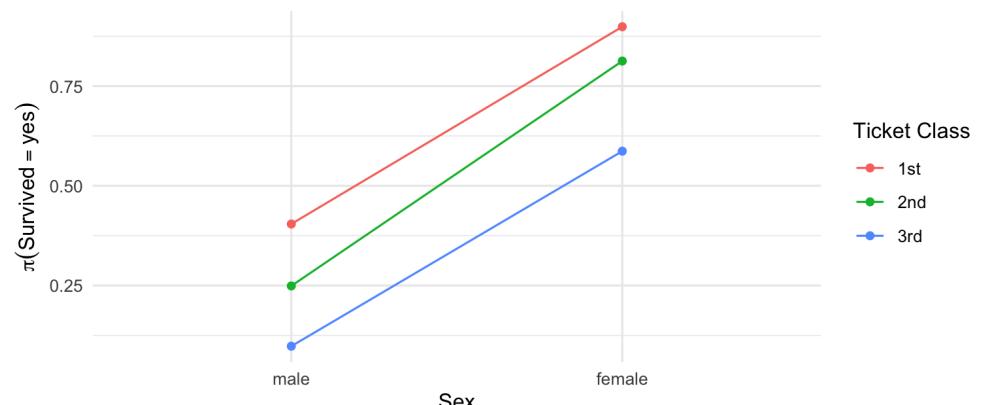
lr_fit1 %>% tidy()
```

```
# A tibble: 4 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>    <dbl>     <dbl>   <dbl>
1 (Intercept)  2.18     0.249     8.77 1.76e-18
2 Sexmale     -2.57     0.208    -12.4  4.56e-35
3 Pclass2nd   -0.716    0.278    -2.58 9.97e- 3
4 Pclass3rd   -1.83     0.247    -7.44 1.03e-13
```

```
newdata <-
  expand.grid(Pclass = c("1st", "2nd", "3rd"),
             Sex = c("male", "female"))

pihat <-
  (lr_fit1 %>% predict(newdata, type = "prob")) %>%
  .pred_Yes

newdata %>% mutate(Pihat = pihat) %>%
  ggplot(aes(Sex, Pihat,
             group = Pclass, colour = Pclass)) +
  geom_line() + geom_point() +
  labs(x = "Sex", y = expression(pi(Survived == yes)),
       colour = "Ticket Class") +
  theme_minimal()
```



Modeling (Titanic case)

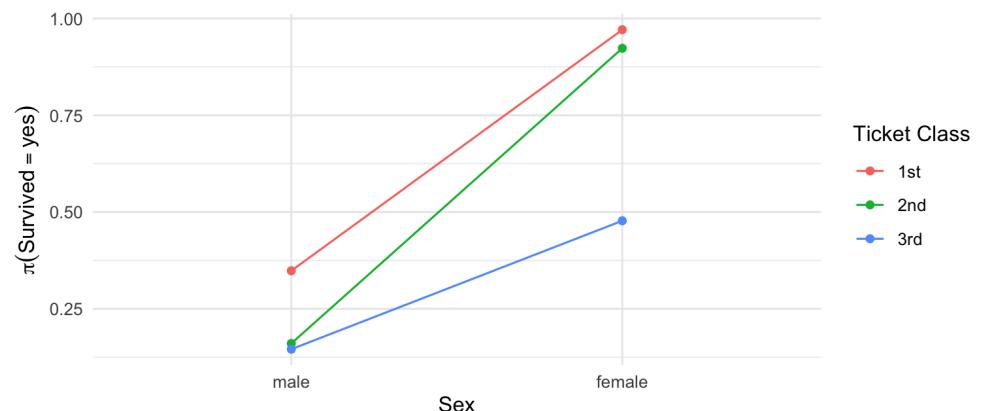
Survived ~ Sex + Pclass + Sex:Pclass

```
lr_fit2 <-
  lr_mod %>%
  fit(Survived ~ Sex * Pclass, data = train_data)

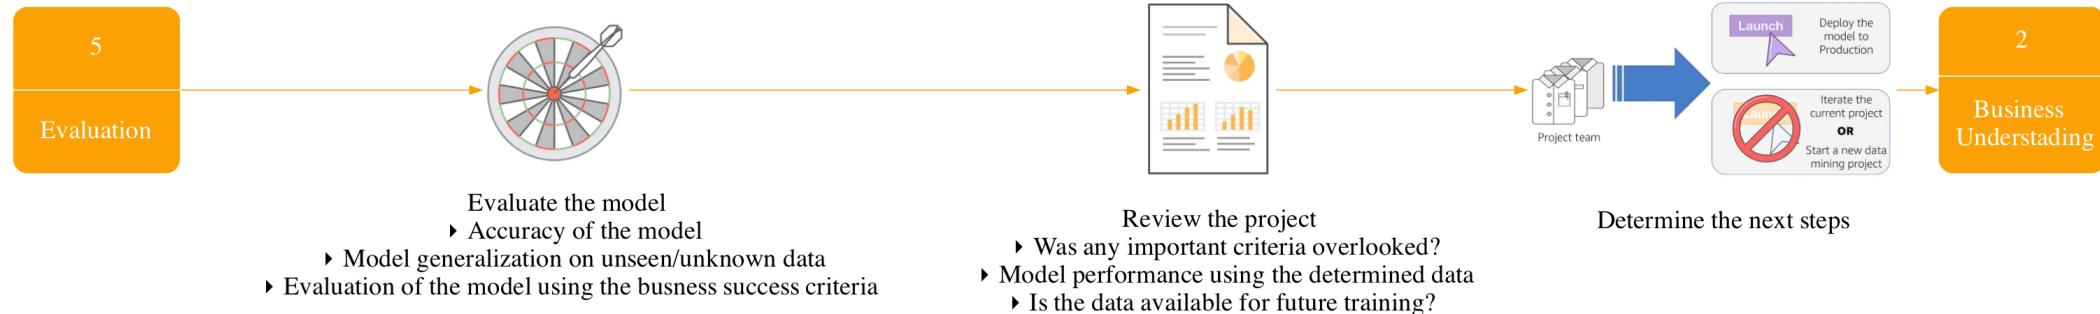
lr_fit2 %>% tidy()
```

| | term | estimate | std.error | statistic | p.value |
|---|-------------------|----------|-----------|-----------|-------------|
| | <chr> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | (Intercept) | 3.51 | 0.718 | 4.89 | 0.000000990 |
| 2 | Sexmale | -4.14 | 0.751 | -5.51 | 0.000000363 |
| 3 | Pclass2nd | -1.03 | 0.855 | -1.20 | 0.230 |
| 4 | Pclass3rd | -3.60 | 0.742 | -4.85 | 0.00000122 |
| 5 | Sexmale:Pclass2nd | -0.00146 | 0.934 | -0.00157 | 0.999 |
| 6 | Sexmale:Pclass3rd | 2.46 | 0.795 | 3.09 | 0.00198 |

```
newdata <-
  expand.grid(Pclass = c("1st", "2nd", "3rd"),
             Sex = c("male", "female"))
pihat <-
  (lr_fit2 %>% predict(newdata, type = "prob")) %%%
  .pred_Yes
newdata %>% mutate(Pihat = pihat) %>%
  ggplot(aes(Sex, Pihat,
             group = Pclass, colour = Pclass)) +
  geom_line() + geom_point() +
  labs(x = "Sex", y = expression(pi(Survived == yes)),
       colour = "Ticket Class") +
  theme_minimal()
```



Evaluation



Mensures

| Predicted\Actual | Survived | No Survived |
|------------------|---------------------|---------------------|
| Survived | True Positive (TP) | False Positive (FP) |
| No Survived | False Negative (FN) | True Negative (TN) |

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$F_{1-score} = \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

Evaluation (Titanic case)

```
measure <- function(data) {  
  data %>% accuracy(truth = Survived, .pred_class) %>%  
    bind_rows(data %>% f_meas(truth = Survived, .pred_class))  
}
```

Survived ~ Sex + Pclass

```
predict(lr_fit1, test_data) %>%  
  bind_cols(predict(lr_fit1,  
                    test_data, type = "prob")) %>%  
  bind_cols(test_data %>% select(Survived)) %>%  
  measure()
```

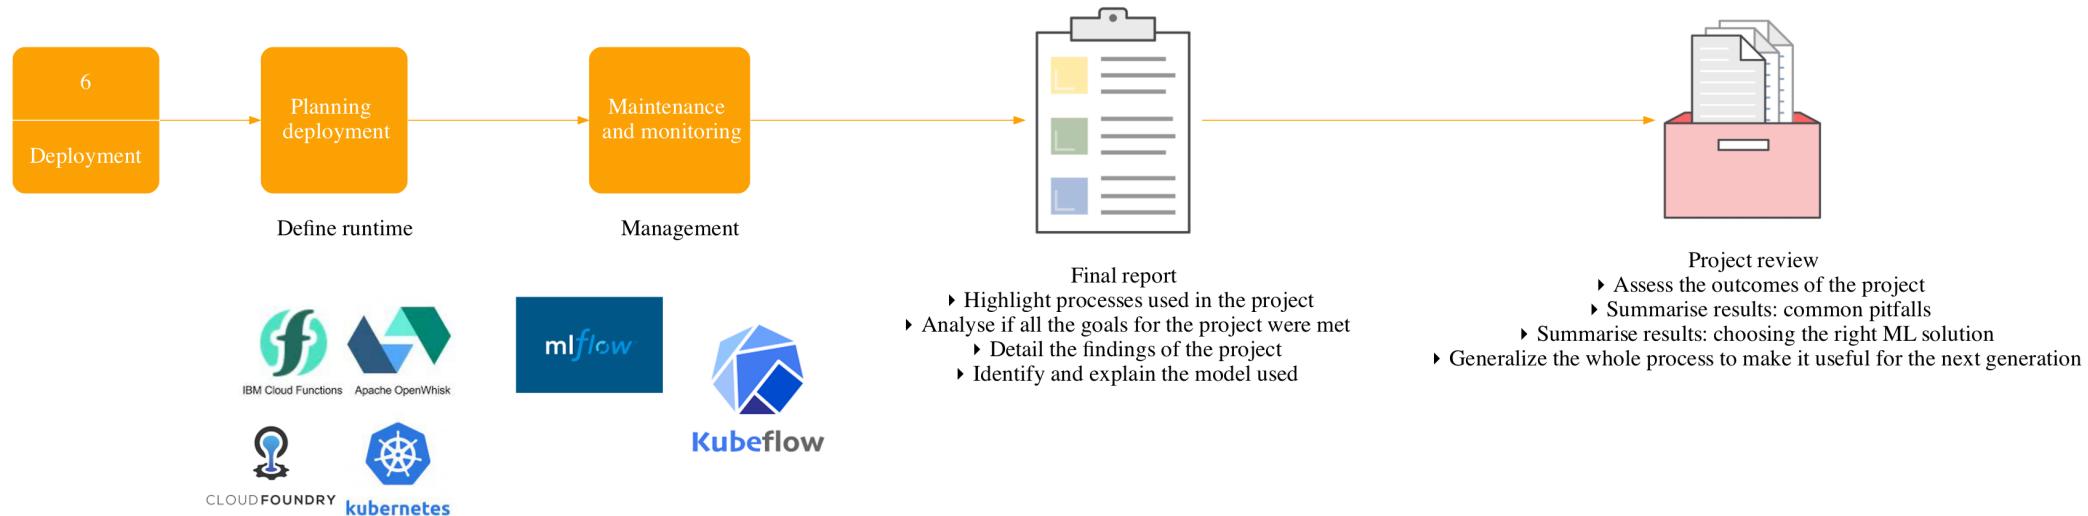
```
# A tibble: 2 x 3  
  .metric  .estimator .estimate  
  <chr>    <chr>        <dbl>  
1 accuracy  binary      0.802  
2 f_meas   binary      0.850
```

Survived ~ Sex + Pclass +
Sex:Pclass

```
predict(lr_fit2, test_data) %>%  
  bind_cols(predict(lr_fit2,  
                    test_data, type = "prob")) %>%  
  bind_cols(test_data %>% select(Survived)) %>%  
  measure()
```

```
# A tibble: 2 x 3  
  .metric  .estimator .estimate  
  <chr>    <chr>        <dbl>  
1 accuracy  binary      0.779  
2 f_meas   binary      0.850
```

Deployment



Deployment (Titanic case)

Save model

```
yaml::write_yaml(  
  tidypredict::parse_model(lr_fit1),  
  "R/my_model.yml"  
)
```

Deployment (Titanic case)

IBM Cloud Functions

IBM Cloud™ Functions service is an event-driven compute platform, also referred to as Serverless computing, or as Function as a Service (FaaS), that runs code in response to events or direct invocations.

Configure

exec

```
#!/bin/bash

# run R script
chmod +x script.R # turn executable
echo "$@" > input.json # set input
./script.R # run script
```

Dockerfile

```
FROM openwhisk/dockerskeleton
RUN apk update && apk add R R-dev R-doc build-base
RUN R -e "install.packages(c('jsonlite', 'tidypredict', 'yaml'), repos = 'http://cran.rstudio.com/')"
```

Deployment (Titanic case)

IBM Cloud Functions

Configure

script.R

```
#!/usr/bin/env Rscript

# load model
loaded_model <-
tidypredict::as_parsed_model(
yaml::read_yaml("my_model.yml"))

# input
input <- jsonlite::fromJSON("input.json", flatten = FALSE)

# compute prediction
pred <- tidypredict::tidypredict_to_column(as.data.frame(input), loaded_model)

# output
jsonlite::stream_out(pred, verbose = FALSE)
```

Deployment (Titanic case)

IBM Cloud Functions

Deploy

```
# docker
docker build th1460/titanic .
docker push th1460/titanic

# login
ibmcloud login -sso
ibmcloud target --cf

# zip
zip -r titanic.zip exec script.R my_model.yml

# deploy
ibmcloud fn action create titanic titanic.zip --docker th1460/titanic --web true
```

Deployment (Titanic case)

IBM Cloud Functions

Request

```
input <- list(Sex = "male", Pclass = "3rd")

"https://us-south.functions.appdomain.cloud/api/v1/web/thiago.pires%40ibm.com_dev/default/titanic.json" %>%
  httr::POST(., body = input, encode = "json") %>%
  httr::content() %>%
  .[c("Sex", "Pclass", "fit")] %>%
  jsonlite::toJSON(pretty = TRUE, auto_unbox = TRUE)

{
  "Sex": "male",
  "Pclass": "3rd",
  "fit": 0.0979
}
```

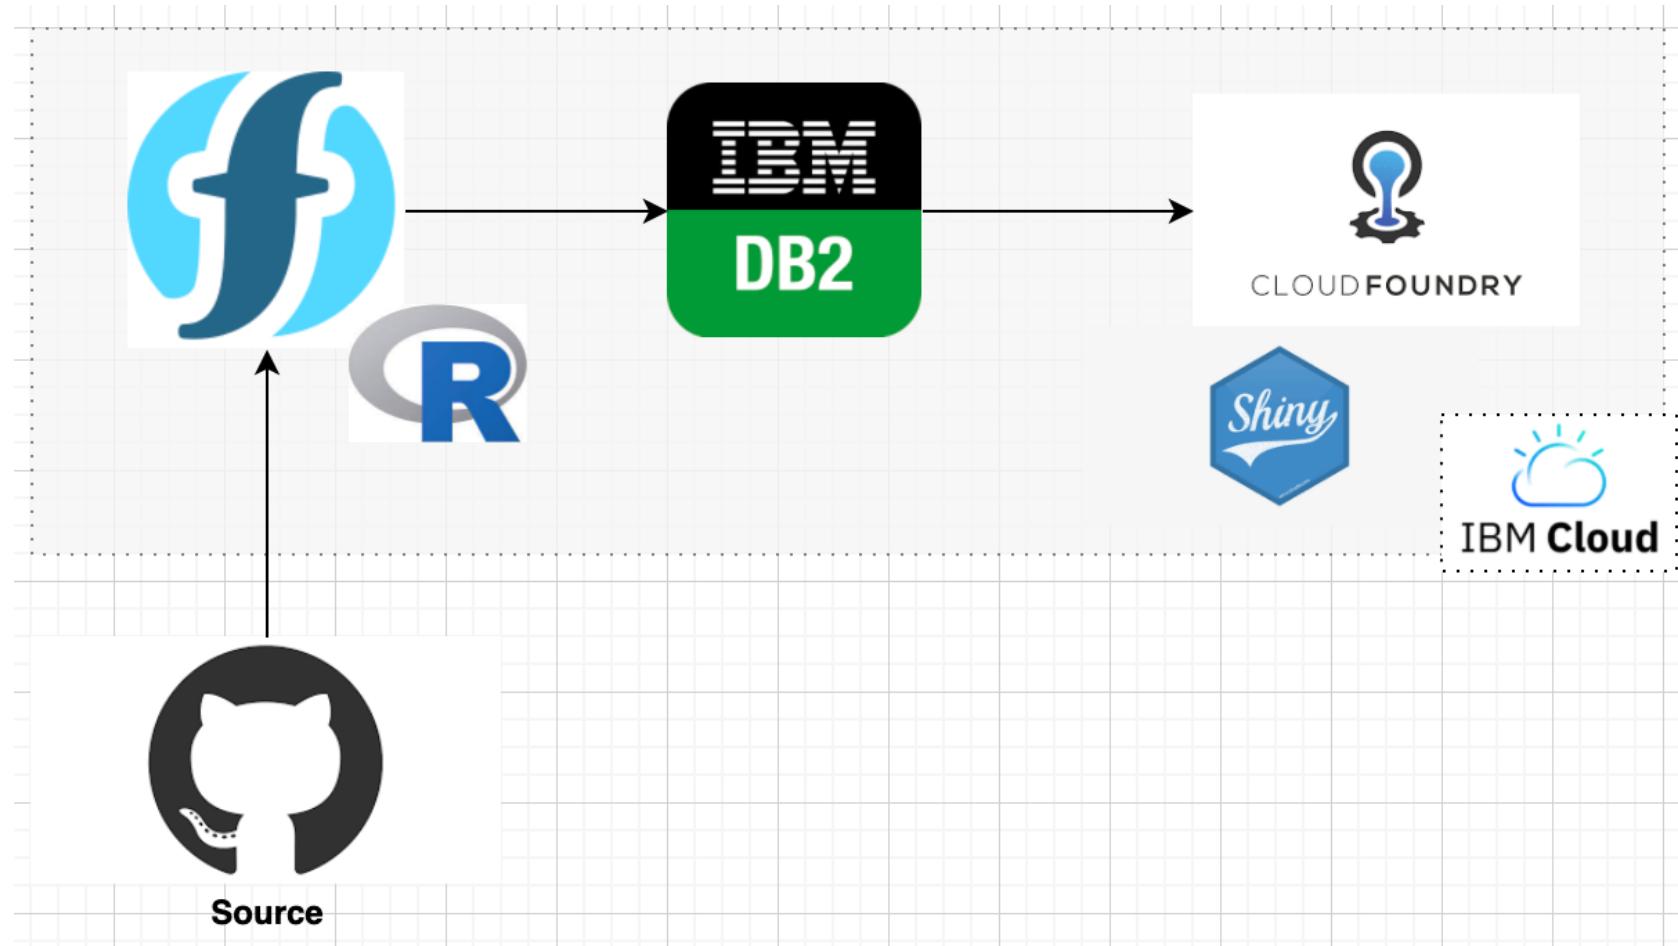
Deployment (Titanic case)

IBM Cloud Functions

The screenshot shows the IBM Cloud Functions deployment interface. On the left, there's a code editor with tabs for "Código" (Code) and "Docker". A message indicates that Docker, Java (.jar), and compacted actions (.zip) cannot be edited in the UI. On the right, there's a section for "Ativações" (Activations) which shows a message: "Nenhuma ativação aqui. Chame sua ação e veja os resultados." (No activation here. Call your action and see the results.)

Covid 19 Analysis

Covid 19 Analysis



Covid 19 Analysis

IBM Functions

- Dockerfile.update
- exec
- read_update_data.R
- .Renvironment

```
# Docker build
docker build -t th1460/update-covid -f Dockerfile.update .
```

```
# Docker push
docker push th1460/update-covid
```

```
# Package files
zip -r update_covid.zip exec read_update_data.R .Renvironment
```

```
# Create function
ibmcloud fn action create update-covid19 update_covid.zip --docker th1460/update-covid --web true
```

Covid 19 Analysis

Cloud foundry

- Dockerfile.app
- app.R
- manifest.yml

```
# Docker build
docker build -t th1460/dash-covid -f Dockerfile.app .

# Docker push
docker push th1460/dash-covid

# Deploy shiny
ibmcloud cf push dash-covid --docker-image th1460/dash-covid -k 2G
```

Covid 19 Analysis

Dashboard



References

- <https://www.aws.training/Details/eLearning?id=27200>
- <https://the-modeling-agency.com/crisp-dm.pdf>



th1460.github.io
github.com/th1460
linkedin.com/in/thop
slack: @thop